

Distributed Hash Tables

Chord

Smruti R. Sarangi

Department of Computer Science
Indian Institute of Technology
New Delhi, India

Outline

- 1 Overview
- 2 Design of Chord
 - Basic Structure
 - Algorithm to find the Successor
 - Node Arrival and Stabilization
- 3 Results

Comparison with Pastry

Chord vs Pastry

- Each node and each key's id is hashed to a unique value.
- The process of lookup tries to find the immediate successor to a key's id.
- The routing table at each node contains $O(\log(n))$ entries.
- Inserting and deleting nodes requires $O(\log(n)^2)$ messages.
- **Sarangi View** 😊 : More robust than Pastry, and more elegant.

Comparison with other Systems

- The **Globe** system assigns objects to locations, and is hierarchical. Chord is completely distributed and decentralized.
- CAN
 - Uses a d -dimensional co-ordinate space.
 - Each node maintains $O(d)$ state, and the lookup cost is $O(dN^{1/d})$.
 - Maintains a lesser amount of state than Chord, but has a higher lookup cost.

Features of Chord

- Automatic load balancing
- Fully distributed
- Scalable in terms of state per node, bandwidth, and lookup time.
- Always available
- Provably correct.


Outline

- 1 Overview
- 2 Design of Chord
 - Basic Structure
 - Algorithm to find the Successor
 - Node Arrival and Stabilization
- 3 Results

Consistent Hashing

Definition

Consistent Hashing: It is a hashing technique that adapts very well to resizing of the hash table. Typically k/n elements need to be reshuffled across buckets. k is the number of keys and n is the number of slots in a hash table.



Structure of Chord

- Each node and key is assigned a m bit identifier.
- The hash for the node and key is generated by using the SHA-1 algorithm.
- The nodes are arranged in a circle (recall Pastry).
- Each key is assigned to the smallest node id that is larger than it. This node is known as the **successor**.

Objective

- For a given key, efficiently locate its successor.
- Efficiently manage addition and deletion of nodes.

Properties of Chord's Hashing Algorithm

- For n nodes, and k keys, with **high probability**
 - ① Each node stores at most $(1 + \epsilon)k/n$ keys
 - ② Addition and deletion of nodes leads to a reshuffling of $O(k/n)$ keys
- Previous papers prove that $\epsilon = O(\log(n))$
- There are techniques to reduce ϵ using virtual nodes.
 - Each node contains $\log(n)$ virtual nodes.
 - Not scalable (**Not necessarily required**)

Chord's Routing(Finger) Table

Let m be the number of bits in an id

- Node n contains m entries in its finger table.
 - successor \rightarrow next node on the identifier circle
 - predecessor \rightarrow node on the identifier circle
- The i^{th} finger contains:
 - $\text{finger}[i].\text{start} = (n + 2^{i-1}) \bmod 2^m, (1 \leq i \leq m)$
 - $\text{finger}[i].\text{end} = (n + 2^i - 1) \bmod 2^m$
 - $\text{finger}[i].\text{node} = \text{successor}(\text{finger}[i].\text{start})$


Basic Operation

$\text{findSuccessor}(\text{keyId}) \rightarrow \text{nodeId}$

Outline

- 1 Overview
- 2 Design of Chord
 - Basic Structure
 - **Algorithm to find the Successor**
 - Node Arrival and Stabilization
- 3 Results

Finger Table- II



Algorithms


Algorithm 1: *findSuccessor* in Chord

```
1 n.findSuccessor(id) begin
2   |   n' ← findPredecessor(id)
   |   return n'.successor(id)
3 end
4 n.findPredecessor(id) begin
5   |   n' ← n
   |   while  $id \notin (n', n'.successor())$  do
6   |   |   n' ← n'.closestPrecedingFinger(id)
7   |   end
8 end
```

closestPrecedingFinger(id)

```
1 n.closestPrecedingFinger(id) begin  
2   for  $i \leftarrow m$  to 1 do  
3     if  $\text{finger}[i].\text{node} \in (n, \text{id})$  then  
4       return  $\text{finger}[i].\text{node}$   
5     end  
6   end  
7   return  $n$   
8 end
```

$O(\log(n))$ Routing Complexity



Outline

- 1 Overview
- 2 Design of Chord
 - Basic Structure
 - Algorithm to find the Successor
 - Node Arrival and Stabilization
- 3 Results

Node Arrival

Each node maintains a predecessor pointer

- Initialize the predecessor and the fingers of the new node.
- Update the predecessor and fingers of other nodes
- Notify software that the node is ready

Node Arrival - II

n initially contacts *n'*

```
1 n.join(n') begin  
2   |   n.initFingerTable(n')  
   |   updateOthers()  
3 end
```

Algorithm 2: *initFingerTable* in Chord

```
1 n.initFingerTable( $n'$ ) begin
2   finger[1].node  $\leftarrow n'$ .findSuccessor(finger[1].start)
   successor  $\leftarrow$  finger[1].node
   predecessor  $\leftarrow$  successor.predecessor
   successor.predecessor  $\leftarrow n$ 
   for  $i \leftarrow 1$  to  $m-1$  do
3     if  $finger[i+1].start \in (n, finger[i].node)$  then
4       finger[i+1].node  $\leftarrow$  finger[i].node
5     end
6     else
7       finger[i+1].node  $\leftarrow n'$ .findSuccessor(finger[i+1].start)
8     end
9   end
10 end
```

updateOthers()

```
1 n.updateOthers() begin
2   for  $i \leftarrow 1$  to  $m$  do
3      $p \leftarrow \text{findPredecessor}(n - 2^{i-1})$ 
4      $p.\text{updateFingerTable}(n, i)$ 
5   end
6 n.updateFingerTable( $s, i$ ) begin
7   if  $s \in (n, \text{finger}[i].\text{node})$  then
8      $\text{finger}[i].\text{node} \leftarrow s$ 
9      $p \leftarrow \text{predecessor}$ 
10     $p.\text{updateFingerTable}(s, i)$ 
11   end
12 end
```

Stabilization of the Netowrk


```
1 n.stabilize() begin
2    $x \leftarrow \text{successor.predecessor}$ 
3   if  $x \in (n, \text{sucessor})$  then
4      $\text{successor} \leftarrow x$ 
5   end
6    $\text{successor.notify}(n)$ 
7 end
8 n.notify( $n'$ ) begin
9   if ( $\text{predecessor is null}$ ) OR ( $n' \in (\text{predecessor}, n)$ ) then
10     $\text{predecessor} \leftarrow n'$ 
11  end
```

Results

Evaluation Setup


- Network consists 10^4 nodes
- Number of keys : 10^5 to 10^6
- Each experiment is repeated 20 times
- The major results are on a Chord protocol simulator

Effect of Virtual Nodes




source [1]

Average Path Length



source [1]

PDF of Path Length



source [1]

Other DHT Systems: Tapestry

Tapestry

- 160 block id, Octal digits
- Routing table like pastry (digit based hypercube)
- Does not have a leaf set or neighborhood table.

Other DHT Systems: Kademlia

Kademlia

- Basis of bit-torrent
- Each node has a 128 bit id
- Each digit contains only 1 bit
- Find the closest node to a key
- Values are stored at several nodes
- Nodes can **cache** the values of popular keys.

Other DHT Systems: CAN

CAN – Content Addressable Network

- It uses a d -dimensional multi-torus as its overlay network.
- Node uses standard routing algorithms for tori. It uses $O(d)$ space. (Note: This is independent of n)
- Each node contains a virtual co-ordinate zone.
- Node Arrival: Split a zone
- Node Departure: Merge a zone



Chord: A Peer-to-Peer Lookup Service for Internet Applications, by I. Stoica, R. Morris, D. Karger, F. Kaashoek, H. Balakrishnan, Proc. ACM SIGCOMM, San Diego, CA, September 2001.